

## Assignment 4: Parallel programming with OpenMP and MPI

Assigned: Thursday, November 8, 2012

Due: 11:55pm Thursday, November 15, 2012

### Introduction

Consider solving the heat diffusion equation

$$\frac{\partial T}{\partial t} = \kappa \nabla^2 T \quad (1)$$

with  $\kappa = \text{constant}$  on a two-dimensional domain of size  $0 \leq x, y \leq \pi$ . Let the boundary conditions be

$$\begin{aligned} T(x, 0) &= \cos^2 x \\ T(x, \pi) &= \sin^2 x \\ T(0, y) &= T(\pi, y) \quad (\text{periodic in } x) \end{aligned} \quad (2)$$

In this assignment, you will write three implementations to solve this problem using finite differences.

1. Write a C, C++ or Fortran program that solves this problem starting from the initial condition  $T = 0$  everywhere. Run the solution on four grids until time  $t = 0.5\pi^2/\kappa$ . Use grid sizes of  $128^2$ ,  $256^2$ ,  $512^2$ , and  $1024^2$ . Plot contours or an image of the final temperature, and report the final, volume averaged temperature.
2. Now extend the program you wrote in part (1) to run with OpenMP. Repeat the calculation using 1, 2, 4, and 8 threads, and report the speed-up you get in each case.
3. Now extend the program you wrote in part (1) to run with MPI. Repeat the calculation using 1, 2, 4, 8, and 16 processes, and again report the speed up. Plot contours of the final temperature and report the volume-averaged value (which will require an MPI\_Allreduce call), and describe how you handled I/O from more than one processor.
4. Discuss the advantages and disadvantages of parallelizing this problem with OpenMP versus MPI.

## Finite difference method

This equation can be solved by centered finite differences in space and the forward Euler method in time,

$$T_{i,j}^{n+1} = T_{i,j}^n + \Delta t \kappa \left( \frac{T_{i-1,j}^n + T_{i+1,j}^n + T_{i,j-1}^n + T_{i,j+1}^n - 4T_{i,j}^n}{\Delta x^2} \right) \quad (3)$$

where  $\Delta x = \Delta y$  is the grid spacing,  $n$  denotes the time step, and  $\Delta t < \frac{\Delta x^2}{4\kappa}$  for numerical stability.

## Implementations

- **Serial:** For the serial version, write a program `heat_serial` that runs with command line options `./heat_serial <nx>` for a solution with grid size  $nx^2$ .
- **OpenMP:** From the serial version, write a parallel version `heat_omp` that runs with command line options `./heat_omp <nx> <nthreads>`. The OpenMP version should be parallelized using the `!$OMP PARALLEL DO` (Fortran) or `#pragma omp parallel for` (C/C++) directive on the appropriate loops.
- **MPI:** From the serial version, write a parallel version `heat_mpi` that runs with `mpiexec ./heat_mpi <nx>`. Parallelize this using *domain decomposition*.

## Domain decomposition

Consider a  $16 \times 16$  grid in figure 1a that we want to run on 4 processors, where the solution variable is in the center of the cells. The easiest way to decompose this is into 4 identical slices of size  $16 \times 4$  as in figure 1b.

Note now instead of solving a single domain of size  $16^2$ , each MPI process solves for  $16 \times 4$  elements, with storage for  $16 \times 6$  elements to account for the *halo* or *ghost* cells indicated in figure 1c. At each time step, you will need to send and receive the appropriate columns of data between processes to populate the ghost cells.

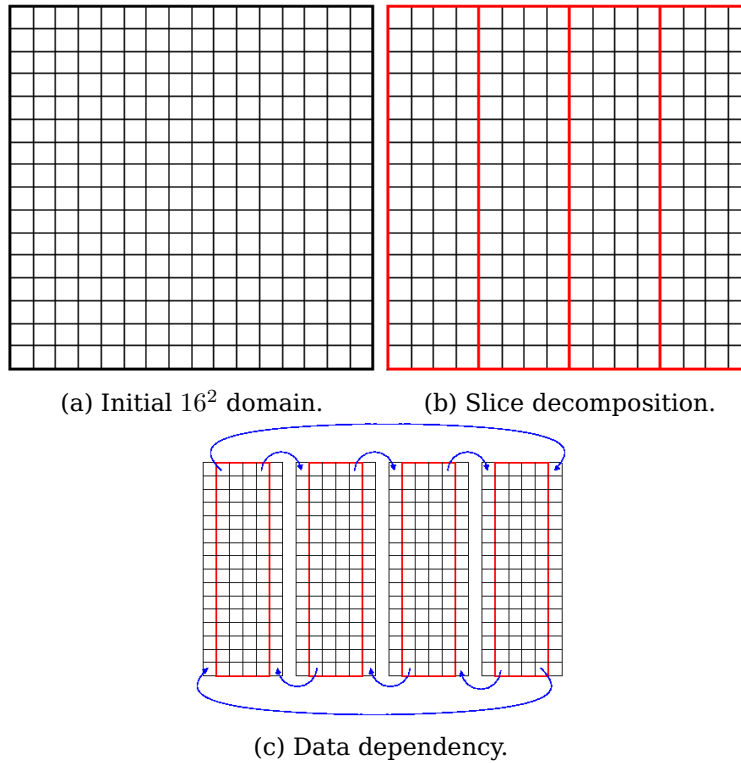


Figure 1: Domain decomposition

## Running parallel jobs

You can develop and test your codes locally, but for the larger tests you should use the Adroit cluster. Information on this cluster is available here: <http://www.princeton.edu/researchcomputing/computational-hardware/adroit/>. You'll need to register for an account if you don't already have one.

Jobs on the Princeton research computing clusters are handled by submitting scripts to the PBS scheduler. For testing the OpenMP and MPI versions on 8 cores, a submission script `heat.run.8` would look like this:

```
#!/bin/bash
#
# parallel/serial job using 1 node and 8 processor cores per node,
# and runs for 1 hour (max).
#PBS -l nodes=1:ppn=8,walltime=01:00:00
#
# send email on abortion, beginning and ending
#PBS -m abe
#
# error and log files
```

```

#PBS -e log.err
#PBS -o log.log
#
#PBS -V

# Load openmpi environment
module load openmpi

cd /home/mlrohry/apc524_hw4/

for nx in 128 256 512 1024
do
    ./heat_omp $nx 8 > heat_omp.$nx.8.out
    # Don't use InfiniBand (IB) on adroit - there's no IB & get warnings without the
    mpiexec --mca btl ^openib ./heat_mpi $nx > heat_mpi.$nx.8.out
done

```

The line `PBS -l nodes=1:ppn=8,walltime=1:00:00` tells the scheduler to reserve 8 processors-per-node on one node, for a maximum of 1 hours. The standard outputs of the programs are being sent to the `.out` files, where the program outputs timing information. You should have a different PBS script for runs with a fixed number of processors; if you reserve 8 processors for runs that only use 1, it will take longer to be scheduled, and prevent other users from accessing the remaining 7 processors.

Adroit has 8 cores per node, so the OpenMP variant can only run on up to 8. For the MPI version using 16 processors, modify the nodes line to `PBS -l nodes=2:ppn=8,walltime=0:30:00` to use 8 cores on each of 2 nodes. Jobs are submitted by the command `qsub heat.run.8`, and can be checked with `qstat -u $USER`.

## Submission

Your submission must include the following:

- `hw4_summary.pdf`: A document containing the relevant plots, and brief discussion of advantages and disadvantages of your OpenMP and MPI implementations.
- Source code for serial, OpenMP, and MPI implementations.
- `Makefile`: Makefile to build the above sources.
- `heat.run.1`, `heat.run.2`, `heat.run.4`, `heat.run.8`, `heat.run.16`: PBS job submission files to run all jobs.

As before, submit these files as a bundled Git repository, `hw4.bundle`. To create a suitable bundle, execute the following from within your repository:

```
git bundle create hw3.bundle master
```

More information on Git can be found at <http://git-scm.com/documentation> and in the slides from lecture.

When you are finished, submit the assignment using the CS Dropbox system at [https://dropbox.cs.princeton.edu/APC524\\_F2012/Parallel](https://dropbox.cs.princeton.edu/APC524_F2012/Parallel)